

# Feature Engineering

Techniques & Use Cases



— *Joy Qi*



# References



Feature Engineering Techniques



Fundamental Techniques of Feature  
Engineering for Machine Learning



`demo.ipynb` in Github Repo

# THE THINKING PROCESS..



# Techniques

## Data Type

- Numerical
- Categorical
- Date/Timestamp
- String

1. Imputation
2. Handling Outliers
3. Binning
4. Grouping
5. Transform
6. Scaling
7. Encoding
8. Extraction

.....

## Data Problem

- Missing values
- Outliers
- Skewness
- Overfitting\*

WHAT

HOW

WHY

# The Menu

Numerical

Categorical

Timestamp



## 1. Imputation

- Just Drop it!
- Numerical Imputation
- Categorical Imputation



## 2. Outliers

- Drop it
- Cap it



## 3. Binning

- Numerical Binning
- Categorical Binning



## 4. Grouping

- Categorical Groupings
- Aggregated sum, mean etc



## 5. Transform

- Log, Square Root
- Box-Cox



## 6. Scaling

- Normalization
- Standardization



## 7. Encoding

- One-Hot Encoding
- Label Encoding



## 8. Extraction

- Date Extraction: Year, Month etc
- IsWeekend, IsHoliday
- Sliding Time Window



# The Menu

Numerical

Categorical

Timestamp



## 1. Imputation

- Just Drop it!
- Numerical Imputation
- Categorical Imputation



## 2. Outliers

- Drop it
- Cap it



## 3. Binning

- Numerical Binning
- Categorical Binning



## 4. Grouping

- Categorical Groupings
- Aggregated sum, mean etc



## 5. Transform

- Log, Square Root
- Box-Cox



## 6. Scaling

- Normalization
- Standardization



## 7. Encoding

- One-Hot Encoding
- Label Encoding



## 8. Extraction

- Date Extraction: Year, Month etc
- IsWeekend, IsHoliday
- Sliding Time Window





# 1. Imputation

**Use Case:** Missing values (NA/NaN/NULL)

**Just drop it!**

**Numerical  
Imputation**

**Categorical  
Imputation**



# 1. Imputation

**Use Case:** Missing values (NA/NaN/NULL)

**Just drop it!**

```
#Dropping rows with NA% higher than threshold  
row_threshold = 0.01  
df = df.loc[df.isnull().mean(axis=1) < row_threshold]
```

```
#Dropping column with NA% higher than threshold  
col_threshold = 0.99  
df = df.loc[df.isnull().mean() < col_threshold]
```

**Note:** Always evaluate the percentage of missing values(NA%) first





# 1. Imputation

**Use Case:** Missing values (NA/NaN/NULL)

**Numerical  
Imputation**

```
#Filling all NAs with certain value
```

```
val = 0
```

```
df_filled = df.fillna(val)
```

```
#Filling NA with median/mean of each column
```

```
df_median = df.fillna(df.median())
```



# 1. Imputation

**Use Case:** Missing values (NA/NaN/NULL)

**Categorical  
Imputation**

```
#Replace all the NULLs of gender column with 'Other'  
df['gender'].replace(np.NaN, 'Other', inplace=True)
```



# 2.Outliers

## Detection & Handling

Use Visualization to detect



Standard Deviation

Percentile

z-score

Drop it

Cap it

### #Dropping the outlier rows with Standard Deviation

```
factor = 3
upper_lim = df['column'].mean() + df['column'].std() * factor
lower_lim = df['column'].mean() - df['column'].std() * factor

df = df[(df['column'] < upper_lim) & (df['column'] > lower_lim)]
```

### #Dropping the outlier rows with Percentiles

```
upper_lim = df['column'].quantile(.95)
lower_lim = df['column'].quantile(.05)

df = df[(df['column'] < upper_lim) & (df['column'] > lower_lim)]
```

### #Capping the outlier rows with Percentiles

```
upper_lim = df['column'].quantile(.95)
lower_lim = df['column'].quantile(.05)

df.loc[(df[column] > upper_lim),column] = upper_lim
df.loc[(df[column] < lower_lim),column] = lower_lim
```



# 3. Binning

**Use Case:** Need for segmentation, Problem of Overfitting

- Continuous to discrete data
- Reduce num. of categories
- Drawback: Sacrifice of information for performance

Numerical Binning

Categorical Binning



## Numerical Binning

Value		Bin
0-30	—>	Low
31-70	—>	Mid
71-100	—>	High

### #Numerical Binning Example

```
df['Bin'] = pd.cut(df['Value'], bins=[0,30,70,100], labels=["Low", "Mid", "High"])
```

	value	bin
0	2	Low
1	45	Mid
2	7	Low
3	85	High
4	28	Low

## Categorical Binning

Country		Continent
Spain	→	Europe
Italy	→	Europe
Brazil	→	South America
Singapore	→	Asia

### #Categorical Binning Example

```
conditions = [df['Country'].str.contains('Spain'),
              df['Country'].str.contains('Italy'),
              df['Country'].str.contains('Brazil'),
              df['Country'].str.contains('Singapore')]

choices = ['Europe', 'Europe', 'South America', 'Asia']

df['Continent'] = np.select(conditions, choices, default='Other')
```



# 4. Grouping

**Use Case:** Transaction datasets where one instance has multiple records hence needs to group by instance and apply aggregate functions

## 1. Categorical Grouping

- Count for frequency
- Pivot Table

	City	Counts
1	London	45678
2	Tokyo	23456
...	...	...

`.value_counts()`

## 2. Numerical Grouping

- Aggregated sum, mean

	month	Sum
1	August	34000
2	September	56000
...	...	...

`.groupby()`



User	City	Visit_Days
1	Roma	1
2	Madrid	2
1	Madrid	1
3	Paris	1
2	Paris	4
1	Paris	3
1	Roma	3



## Pivot Table

*(Similar to One-Hot Encoding)*

User	Paris	Madrid	Roma
1	3	1	4
2	4	2	0
3	1	0	0

### #Pivot table Pandas Example

```
df.pivot_table(index='column_to_group',
                columns='column_to_encode',
                values='aggregation_column',
                aggfunc=np.sum,
                fill_value = 0)
```



# 5.Transform

**Use Case:** Data skewness which violates model assumptions

## 1. Log Transformation:

- **Typical example:** the vast majority of incomes are small and very few are big
- **Normalize the magnitude difference:** skewed to normal distribution
- **Decrease outliers effect:** model become more robust

**Note:** Can only be applied to positive values, negative value will return NULL

```
#Log Transform Example with Negative values
df = pd.DataFrame({'value':[2,45, -23, 85, 28, 2, 35, -12]})

#Negative Values Handling
#Choose constant c where  $c=1-X_{\min}$  so that  $\min(X+c)=1$  and  $(X-X_{\min}+1)$  will always  $\geq 1$  ( $\log(1)=0$ )
df['log'] = (df['value']- df['value'].min()+1).transform(np.log)
```

## 2. Other Transformations: Square Root, Box-Cox



# 6. Scaling

**Use Case:** Compare columns with different ranges (e.g. Age and income)

1. **Normalization: Scale all column values to same fixed range [0, 1]**

$$X_{\text{norm}} = (X - X_{\text{min}}) / (X_{\text{max}} - X_{\text{min}})$$

```
df = pd.DataFrame({'value': [2, 45, -23, 85, 28, 2, 35, -12]})  
df['normalized'] = (df['value'] - df['value'].min()) / (df['value'].max() - df['value'].min())
```

2. **Standardization: Take consideration of std to get different ranges**

$$z = (X - X_{\text{mean}}) / \text{std}$$

```
df = pd.DataFrame({'value': [2, 45, -23, 85, 28, 2, 35, -12]})  
df['standardized'] = (df['value'] - df['value'].mean()) / df['value'].std()
```



# 7.Encoding

## 1. Label Encoding

- **Use Case:** Convert categories in 1 column into numerical values
- **Drawback:** Create unnecessary weightage for nominal categories
  - ✔ Apply only to ordinal categories or use One-Hot Encoding

## 2. One-Hot Encoding

- **Use Case:** Convert categories in 1 column into multiple binary columns
- **Drawback:** Create too many binary features, data becomes sparse
  - ✔ Selectively apply to the most frequent categories

Std	Grade
1	A
2	D
3	B
4	A
5	C
6	D

Ordinal  
Categories



Label  
Encoding

Std	Grade	Value
1	A	5
2	D	2
3	B	4
4	A	5
5	C	3
6	D	2

Id	Season
1	Spring
2	Winter
3	Summer

Nominal  
Categories



Id	Season	Value
1	Spring	1
2	Winter	4
3	Summer	3

User	City
1	Tokyo
2	Sydney
3	Tokyo
4	London
5	London
6	Honolulu



## One-Hot Encoding

User	Tokyo	Sydney	London	Honolulu
1	1	0	0	0
2	0	1	0	0
3	1	0	0	0
4	0	0	1	0
5	0	0	1	0
6	0	0	0	1

	City	Counts
1	London	45678
2	Tokyo	23456
...	...	...
512	Honolulu	1
513	Budapest	1
514	Bangkok	1

Apply to the most frequent ones

Park under 'Others'



User	London	Tokyo	Others
1	1	0	0
2	0	1	0
...	...	...	...

## Label Encoding for the Grade

```
#Create map of category:value
cat_dict = {'A': 5, 'B': 4, 'C': 3, 'D': 2}
df['Value']=df['Grade'].replace(cat_dict, inplace=False)
```

## One-Hot Encoding for the first 30 most frequent categories

```
#Create category list (30 + 1 Other)
cat_list = df['col'].value_counts()[:30].keys().tolist()
cat_list.append('other')

#Map original column into a new column with only 31 categories
df['new_col']=pd.Categorical(df['col'], categories=cat_list)
df['new_col'].fillna('other', inplace=True)

#One-Hot Encoding
col_encoding = pd.get_dummies(df['new_col'])

#Join the encoding columns back to original DataFrame
df = pd.concat([df, col_encoding], axis=1)
df.drop(['new_col'], axis=1, inplace=True)
```



# 8. Extraction on Date

1. **Timestamp to year, month, day-of-week, hour, minute etc.**
  - **Use Case:** For aggregating features (e.g. Sum per month)
2. **IsWeekday, IsWeekend, IsHoliday**
  - **Use Case:** User behaviour driven (e.g. higher # of shopping transactions during Xmas holidays)
3. **Sliding Time Window**
  - **Use Case:** Aggregating under Dynamic/Customized time window



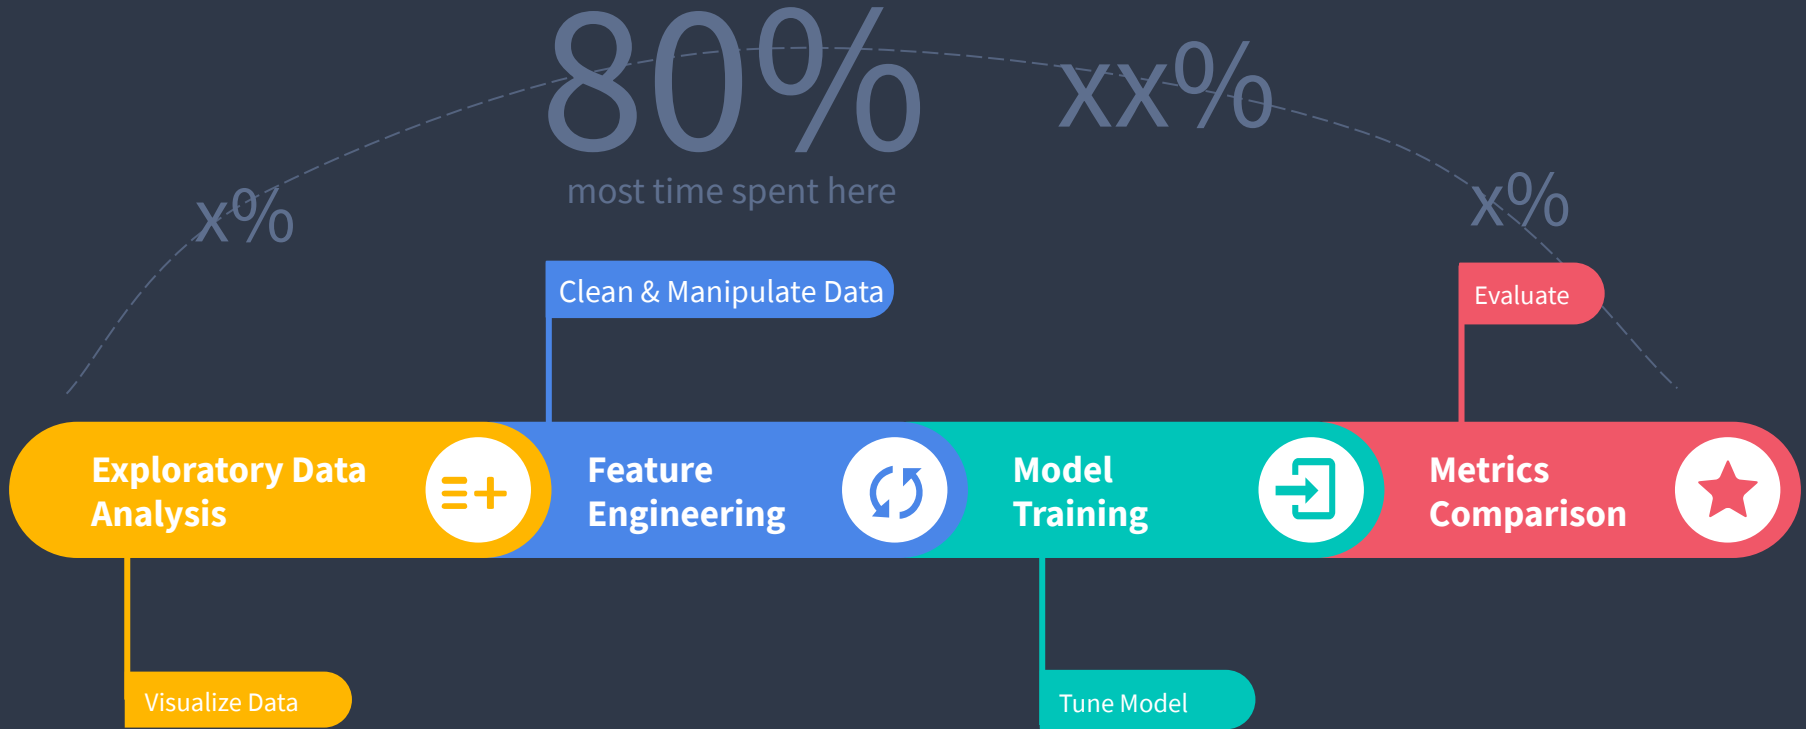
## Extract Date Features: weekday, isweekend, isholiday

```
holiday_list = ['2019-01-01', '2019-01-21', '2019-05-27', '2019-07-04',  
                '2019-09-02', '2019-11-28', '2019-11-29', '2019-12-25']
```

```
def get_date_features(df):  
    """  
    Given a DataFrame with Date columns  
    Create Binary Columns on whether the date:  
    - is weekend: 1, otherwise: 0  
    - is Federal holiday: 1, otherwise: 0  
  
    Return a DataFrame with features created  
    """  
    df['departure_weekday'] = pd.to_datetime(df['departure_date']).dt.weekday  
    df['departure_isweekend'] = (df['departure_weekday'].isin([6, 7])).astype(int)  
    df['departure_isholiday'] = (df['departure_weekday'].isin(holiday_list)).astype(int)  
  
    return df
```

```
df = get_date_features(df)
```

# GENERAL DATA SCIENCE STEPS



# THE TAKEAWAYS

1.



2.

**Feature Engineering** comes from understanding the **Data**

# The Specials for next time



## NLP

- TF-IDF
- StopWords removal
- Continuous Bag-of-Words
- Word Embeddings



## Web Scraping

- BeautifulSoup BS4
- Selenium
- Handling JSON data



## Regularization

- L1: Lasso
- L2: Ridge
- Elastic-net (combination of L1 & L2)



## PCA

- Principal Component Analysis
- linear dimensionality reduction



## K-means

- Elbow Method
- Calculate Distances: Euclidean, Haversin etc.



## Imbalanced data

- Up-sampling
- Down-sampling

# THANK YOU

**Happy Data Science** 🤖🤖🤖🤖

Q&A